

Transformers

Zeham Management Technologies BootCamp
by SDAIA

September 12th, 2024



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority



Agenda



Fine-tuning



Advanced Technologies



Fine-Tuning



Fine-Tuning Transformers

Use pre-trained models as a starting point for custom tasks, reducing training time and resources.

Fine-Tuning Workflow:

1. Load pretrained model.
2. Prepare dataset (tokenization, formatting).
3. Use Hugging Face's Trainer API for training.



Hyperparameter Tuning in Fine-Tuning

Fine-tuning requires careful selection of hyperparameters such as learning rate, batch size, and the number of training epochs.

- **Learning Rate:** Adjusting learning rates can prevent overfitting or underfitting.
- **Batch Size:** Smaller batch sizes give more granular updates but take longer, whereas larger batches improve training speed but may generalize less well.
- **Number of Epochs:** This determines how many times the model will pass through the entire dataset during training.



Learning Rate Scheduling

Transformers have many layers and parameters, making them sensitive to learning rate. Adjusting it dynamically during fine-tuning can lead to better convergence and stability.

Popular Learning Rate Schedulers for Transformers:

- **Warmup with Linear Decay:** Often used in transformers, this technique increases the learning rate gradually at the start (warmup) and then decreases it linearly. Effective for large pre-trained models.
- **Cosine Annealing:** Reduces the learning rate in a cyclic manner, useful for longer fine-tuning tasks where the model requires smooth convergence.



Using Dropout for Transformer Regularization

What is Dropout in Transformers?

Dropout is a regularization technique where some neurons are “dropped out” or turned off during each forward pass to prevent overfitting.

In transformers, dropout can be applied to attention layers and feed-forward layers to ensure that the model generalizes well to unseen data during fine-tuning.

When to Use Dropout in Fine-Tuning:

- Dropout is especially useful when fine-tuning on smaller datasets where overfitting is a concern.
- Applying it to the attention weights ensures that the model doesn't rely too heavily on specific attention patterns.



Layer Freezing for Efficient Fine-Tuning

Why Freeze Layers in Transformers?

Freezing the lower layers of a transformer model can drastically speed up fine-tuning, especially when those layers capture general features from the pretraining process.

By freezing lower layers and only updating the top layers, you reduce computation while focusing learning on task-specific patterns.

When to Use Layer Freezing:

Effective when transferring a transformer model to a new task with limited data, where task-specific features are more important than general language understanding.

Often, only the last few layers are fine-tuned to adapt to the new task, especially for BERT-like models.

Layer Freezing for Efficient Fine-Tuning

Freeze Layers in Transformers Example:

Code Example (Freezing Layers in Hugging Face Transformers):

python

 Copy code

```
# Freeze all layers except the last
for param in model.bert.encoder.layer[:10].parameters():
    param.requires_grad = False
```



Handling Imbalanced Datasets in Fine-Tuning

Why Imbalanced Data Matters:

Transformers can overfit to majority classes when fine-tuning on imbalanced datasets.

Addressing this improves model generalization to minority classes.

Strategies:

- **Oversampling Minority Class:** Duplicate samples from the minority class to balance the dataset.
- **Class Weights:** Assign higher loss weights to the minority class to balance the contribution to the loss function.




Handling Imbalanced Datasets in Fine-Tuning

Class Weights Example:

Code Example (Using Class Weights in Transformers):

python

 Copy code

```
from torch.nn import CrossEntropyLoss

# Define weights for imbalanced classes
class_weights = torch.tensor([0.7, 1.3])

# Use class weights in the loss function
loss_fct = CrossEntropyLoss(weight=class_weights)
```

Advanced Technologies




Advanced Optimization Algorithms

- **AdamW:** Combines adaptive learning rates with weight decay, commonly used in fine-tuning.
- **RMSprop:** Effective when training on highly noisy or non-stationary data.
- **SGD with Momentum:** A variant of stochastic gradient descent that accelerates training by keeping a "momentum" term.

Advanced Optimization Algorithms

AdamW Example:

python

 Copy code

```
from transformers import AdamW

# Define the optimizer using AdamW
optimizer = AdamW(model.parameters(), lr=5e-5, weight_decay=0.01)

# Training loop
for batch in train_dataloader:
    outputs = model(**batch)
    loss = outputs.loss
    loss.backward()

    optimizer.step()
    optimizer.zero_grad()
```



Mixed Precision Training for Faster Fine-Tuning

What is Mixed Precision Training?

Mixed precision involves using both 16-bit and 32-bit floating point types during training to speed up computation and reduce memory usage without sacrificing much accuracy.

Advantages for Transformers:

- Faster training time with less memory usage, especially useful for large models like BERT, GPT-2, or T5.
- Hugely beneficial when training on limited hardware resources.



Mixed Precision Training for Faster Fine-Tuning

Mixed Precision Example:

Code Example (Using Hugging Face's Trainer API with Mixed Precision):

python

 Copy code

```
from transformers import Trainer, TrainingArguments

training_args = TrainingArguments(
    fp16=True, # Enable mixed precision training
    per_device_train_batch_size=16,
    ...
)

trainer = Trainer(
    model=model,
    args=training_args,
    ...
)
```


Tutorial

8-Transformers/LAB/ Question Answering with Transformers.ipynb

Tutorial

8-Transformers/LAB/ Text-to-Text Generation.ipynb

Tutorial

8-Transformers/LAB/ Fill Mask.ipynb

Thank you!



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority